

# Attention Mechanisms in Transformer Architectures: A Brief Survey

幸川卓実

August 29, 2026

## Abstract

This survey examines the role of attention as the central mechanism that enables Transformers to model dependencies across tokens, positions, and modalities. By assigning content-dependent weights to input representations, attention allows the model to selectively aggregate information and capture both local interactions and long-range context, making it the key driver of Transformer performance.

The survey covers the main attention formulations used in Transformer architectures, including scaled dot-product attention, multi-head attention, self-attention, and cross-attention. It also discusses sparse and other efficient variants designed to reduce the quadratic cost of full attention while preserving expressive power. Across these families, we compare their computational trade-offs, representational strengths, and typical use cases in language, vision, and multimodal settings.

Our takeaway is that attention is both the defining strength and the primary bottleneck of Transformers: it provides flexible global context modeling, but its cost motivates continual innovation in approximation, sparsity, and architectural design. Understanding these mechanisms is essential for interpreting current Transformer models and for developing more scalable future systems.

## 1 Introduction

Transformers have become the dominant architecture for sequence modeling because they replace recurrence with attention [5], a mechanism that directly compares all elements in a sequence and learns how much each element should influence every other element. This design makes it possible to model long-range dependencies more effectively than many earlier neural approaches, while also enabling highly parallel computation during training.

The purpose of this paper is to provide a compact survey of the attention mechanisms that make Transformer models work. Rather than focusing on a single application, we emphasize the common mathematical structure shared by attention variants and explain how different design choices affect expressiveness, efficiency, and scalability. In particular, we highlight the trade-off between global context modeling and computational cost, which has driven much of the recent research in this area.

The remainder of the paper is organized as follows. Section 2 introduces the notation and the sequence-to-sequence setting. Section 3 presents scaled dot-product attention and its complexity. Section 4 discusses multi-head attention, while Section 5 contrasts self-attention and cross-attention. Section 6 reviews sparse and efficient variants, and the final sections summarize the main lessons and conclude the survey.

## 2 Preliminaries

### 2.1 Notation

Attention-based models operate on sequences of vectors. We write an input sequence as  $X = (x_1, x_2, \dots, x_n)$ , where each  $x_i \in \mathbb{R}^d$  is a token representation. In practice, these vectors may come from learned embeddings, image patches, audio frames, or other modality-specific encoders. The goal is to transform the input into contextualized representations that incorporate information from the entire sequence.

Throughout this paper, we use  $n$  for sequence length,  $d$  for the model dimension, and  $d_k$  and  $d_v$  for the dimensions of keys and values. We also use matrices  $Q$ ,  $K$ , and  $V$  to denote collections of queries, keys, and values, respectively. This notation is standard in the Transformer literature and makes it easy to express attention as a matrix operation. Sequence-to-sequence learning maps an input sequence to an output sequence, often with different lengths. Classical encoder-decoder systems compress the input into a fixed-size state and then generate outputs autoregressively. While effective, this bottleneck can make it difficult to preserve detailed information from long inputs.

Attention addresses this limitation by allowing the decoder to access all encoder states directly. Instead of relying on a single summary vector, the model computes a weighted combination of source representations at each decoding step. This mechanism improves alignment between source and target elements and provides a more flexible way to handle variable-length dependencies.

## 3 Scaled Dot-Product Attention

### 3.1 Formulation

Scaled dot-product attention computes similarity scores between a query and a set of keys, then uses those scores to form a weighted sum of the corresponding values. Given query matrix  $Q \in \mathbb{R}^{n_q \times d_k}$ , key matrix  $K \in \mathbb{R}^{n_k \times d_k}$ , and value matrix  $V \in \mathbb{R}^{n_k \times d_v}$ , the attention output is

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V. \quad (1)$$

As shown in Equation (1), the scaling factor  $\sqrt{d_k}$  stabilizes optimization by preventing the dot products from growing too large as the dimensionality increases. The dominant cost of full attention comes from forming the  $n_q \times n_k$  similarity matrix, which requires  $O(n_q n_k d_k)$  operations, and from storing that matrix during training. In the common self-attention case where  $n_q = n_k = n$ , this becomes quadratic in sequence length, both in time and memory.

This quadratic scaling is acceptable for moderate sequence lengths, but it becomes a major limitation for long documents, high-resolution images, and other settings with many tokens. The search for more efficient attention mechanisms is therefore largely a search for ways to approximate or restructure this computation without losing the ability to model rich dependencies.

## 4 Multi-Head Attention

Multi-head attention extends the basic attention mechanism by running several attention operations in parallel, each with its own learned projections of the input. For head  $h$ , the model computes

$$\text{head}_h = \text{Attention}(QW_h^Q, KW_h^K, VW_h^V),$$

where  $W_h^Q$ ,  $W_h^K$ , and  $W_h^V$  are learned matrices. The outputs of all heads are concatenated and projected again to produce the final representation.

This design lets the model attend to different types of relationships simultaneously. One head may focus on local syntactic structure, another on long-range semantic dependencies, and another on positional or modality-specific patterns. As a result, multi-head attention increases representational capacity without changing the basic attention formula.

## 5 Self-Attention and Cross-Attention

Self-attention is the special case in which queries, keys, and values all come from the same sequence. Every token can therefore attend to every other token, including itself, which makes the representation of each position context-dependent. This is the mechanism that allows Transformers to build deep interactions within a single sequence, as demonstrated in BERT [3].

Cross-attention, by contrast, uses queries from one sequence and keys and values from another. In encoder-decoder architectures, the decoder queries the encoder states to retrieve source information relevant to the current generation step. This separation is especially useful in translation, summarization, and multimodal fusion, where one sequence conditions another.

## 6 Sparse and Efficient Variants

### 6.1 Sparse Attention

Sparse attention reduces cost by restricting each query to a subset of keys rather than attending to the full sequence [1]. The subset may be chosen using fixed patterns such as local windows, strided connections, or block structures, or it may be learned dynamically. By limiting the number of pairwise interactions, sparse methods can substantially reduce memory use while still preserving useful long-range paths through the network.

Linear attention takes a different approach by rewriting the attention computation so that it avoids explicitly forming the full similarity matrix [2, 6]. These methods typically approximate the softmax kernel or use feature maps that make the computation associative, reducing the complexity from quadratic to linear in sequence length under suitable assumptions. Although such methods may trade away some exactness, they are attractive for very long inputs where standard attention becomes impractical.

## 7 Discussion

The attention family is best understood as a spectrum of design choices around the same core idea: compute relevance scores, normalize them, and use them to mix information. Full attention offers the most direct and expressive version of this idea, but it is also the most expensive. Sparse and linear variants reduce cost, yet they often introduce approximations or structural constraints that can affect accuracy.

In practice, the best choice depends on the task. Language modeling often benefits from dense self-attention at moderate sequence lengths, while long-context applications and high-resolution perception tasks may require sparse or approximate methods. This diversity suggests that attention is not a single algorithm but a flexible framework that can be adapted to different computational budgets and data modalities.

A useful way to think about attention is as a learned routing mechanism. Rather than forcing every token to contribute equally, the model assigns different weights depending on the current context. This makes the representation adaptive: the same word, patch, or frame can play different roles depending on what surrounds it. That flexibility is one reason attention has been so successful across tasks.

At the same time, attention is not a complete solution to sequence modeling. Its quadratic cost can become prohibitive, and its behavior can be difficult to interpret at scale because

attention weights do not always correspond to human-readable explanations. For this reason, recent work has explored hybrids that combine attention with recurrence, convolution, state-space models, or retrieval mechanisms. These alternatives aim to preserve the strengths of attention while reducing its computational burden or improving inductive bias.

Another important point is that attention is highly modular. The same basic operator can be used in encoders, decoders, multimodal fusion blocks, and even non-sequential settings such as graphs or sets. This modularity helps explain why the Transformer template has spread so widely: once the attention primitive is available, it can be inserted into many architectures with relatively small changes.

Overall, the story of attention is one of balance. Dense attention is elegant and powerful, sparse attention is efficient and scalable, and approximate attention is practical for long contexts. The ongoing challenge is to combine these advantages without losing the modeling capacity that made Transformers successful in the first place.

## 8 Conclusion

$$f(x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} F(k) e^{ikx} dk$$

Attention has reshaped modern machine learning by providing a simple yet powerful way to route information through a model. In Transformers, it serves as the core operation that enables contextual representation learning, flexible alignment, and scalable parallel computation.

This survey has reviewed the main attention mechanisms used in Transformer architectures and highlighted the trade-offs among accuracy, interpretability, and efficiency. Future progress will likely come from better sparse patterns, more accurate linear approximations, and hybrid architectures that combine attention with other inductive biases. The success of vision Transformers [4] further demonstrates the generality of these mechanisms beyond language.

## References

- [1] Iz Beltagy, Matthew E Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- [2] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, David Belanger, Lucy Colwell, and Adrian Weller. Rethinking attention with performers. In *International Conference on Learning Representations*, 2021.
- [3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2019.
- [4] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2021.
- [5] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.

- [6] Sinong Wang, Belinda Z Li, Madian Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention with linear complexity. In *International Conference on Learning Representations*, 2020.